

1 Introduction

In my previous paper, I built a simple Bayesian game with a single consumer. This paper extends that setting to a small niche market with heterogeneous consumers each with a different willingness to pay. The producer does not make any parametric assumption about the distribution of consumers' willingness to pay (except regularity conditions, e.g. an atomless F with a density), but only seeks the optimal quantile that maximise the profits. This yields a general Bayesian game well-suited to a DQN. While a fully non-parametric approach can be intractable when maintaining the entire posterior over distribution, I argue that two statistics - the optimal price and the local hazard rate - characterise the producer's learning, making the model tractable. As before, obfuscation introduces strategic noise in demand realisation that attenuates the mapping from willingness-to-pay to observed acceptance rates, thereby distorting the producer's inference of the hazard rate.

2 Theoretical Base

In this game, we have N heterogeneous consumers, each with private willingness to pay W_i . Each consumer's policy is mixed :

$$\epsilon_i \mathbb{1}(W_i \geq x_t) + \frac{(1 - \epsilon_i)}{2}, \quad \epsilon_i \in [0, 1]$$

where ϵ_i controls the level of obfuscation. Consumers try to learn optimal ϵ_i^* via DQN.

The producer is single and follows a belief-greedy strategy: it maintains an estimate $\hat{A}_t(\cdot)$ of the acceptance curve $\tilde{S}(\cdot; \bar{\epsilon})$ updated online from observed acceptance rates, and prices greedily: $x_t = \arg\max_x (x - c) \hat{A}_t(x)$. By the theorem in section 2.1, the producer's optimal price depends on $\hat{A}_t$ only through the two local statistics $(x_t, \hat{h}_t)$ - the current price and the local effective hazard- which we therefore take as its state. The motivation for this dimension reduction is statistical : with N consumers the empirical CDF satisfies $\sup_x |\hat{F}_N(x) - F(x)| = O_p(N^{-1/2})$ (Dvoretzky-Kiefer-Wolfowitz), so under the atomless assumption, the plug-in estimates $(\hat{x}_t, \hat{h}_t)$ are consistent, so a market with sufficiently large N lets the producer recover the two-dimensional statistic with high probability. The producer's expected profit at price x , given aggregate obfuscation $\bar{\epsilon}$, is

$$\Pi(x; \bar{\epsilon}) = (x - c) \tilde{S}(x; \bar{\epsilon}), \quad \tilde{S}(x; \bar{\epsilon}) := \bar{\epsilon}(1 - F(x)) + \frac{1 - \bar{\epsilon}}{2}$$

where c is the unit production cost, $\tilde{S}$ is the effective demand. The constant term $\frac{1 - \bar{\epsilon}}{2}$ is a price insensitive noise floor of acceptances produced by randomisation.

2.1 Reduction of Dimensionality

Lemma 2.1. *Assume F is differentiable with density f and $\tilde{f}(x) = \bar{\epsilon}f(x) > 0$ at any interior maximiser. Then every interior optimal price x^* satisfies*

$$x^* = c + \frac{\tilde{S}(x^*; \bar{\epsilon})}{\tilde{f}(x^*)} = c + \frac{1}{\tilde{h}(x^*; \bar{\epsilon})}$$

pf. $\Pi(x; \bar{\epsilon}) = (x - c)\tilde{S}(x; \bar{\epsilon})$ is differentiable with $\Pi'(x) = \tilde{S}(x) - (x - c)\tilde{f}(x)$. At an interior maximiser $\Pi'(x^*) = 0$; dividing by $\tilde{f}(x^*) > 0$ gives $x^* - c = \tilde{S}(x^*)/\tilde{f}(x^*) = 1/\tilde{h}(x^*)$. Since $\tilde{S}(x; \bar{\epsilon}) \rightarrow \frac{1-\bar{\epsilon}}{2} > 0$ as x grows, the noise floor can push the maximiser to the upper boundary of the price domain; an interior optimum need not exist for $\bar{\epsilon} < 1$. This distortion relative to the truthful benchmark ($\bar{\epsilon} = 1$) is the mechanism by which obfuscation operates. We therefore work on a bounded price domain $x \in [0, \bar{x}]$ on which the lemma characterises any interior optimum and Kakutani gives existence.

Theorem 2.1 (Sufficient Statistic). *The producer's local optimal pricing depends on F only through the local hazard function near x^* .*

pf. By the previous lemma, any interior optimum solves $x = c + 1/\tilde{h}(x; \bar{\epsilon})$, which involves F only through $\tilde{h}(x; \bar{\epsilon}) = \bar{\epsilon}f(x)/\tilde{S}(x; \bar{\epsilon})$ evaluated near x . Hence two type distributions inducing the same effective hazard on a neighbourhood of their optimum share the same interior optimal price; the optimality check depends on F only through the local pair $(x, \tilde{h}(x; \bar{\epsilon}))$.

→ the producer's learning state can be taken to be the two-dimensional $(x_t, \tilde{h}_t)$ which we use as the input to the DQN, paralleling the (μ, σ) state of the single-type model.

2.2 Existence of Equilibrium in the Game

Assume F , the CDF of consumer willingness-to-pay, is atomless. Then the expected demand $\tilde{S}$ is a continuous function of $(x, \bar{\epsilon})$ even though individual accept indicators jump - the discontinuity gets integrated out over the type distribution F . By Berge's maximum theorem, the producer's best-response correspondence $X^*(\bar{\epsilon}) = \arg\max_x \Pi(x; \bar{\epsilon})$ is nonempty-valued and u.h.c.; the induced aggregate best response map on $\bar{\epsilon} \in [0, 1]$ is u.h.c. with nonempty, convex (from mixed obfuscation strategies), compact values. Therefore by the Kakutani's fixed point theorem, the fixed point (Bayes-Nash equilibrium) exists.

→ DQN dynamics approximate this adjustment; the experiment tests whether they converge to the fixed point promised.

3 Empirical Results

The producer observes the consumers’ actions, updates its belief about the acceptance curve, and reprices accordingly. Each consumers maximises its own discounted surplus, knowing producer adapts its price to aggregate behaviour. In this game, we vary β , the time discounting factor, to test whether patience correlates with obfuscation. As β increases, the equilibrium $\bar{\epsilon}^*$ decreases - more

Table 1: Equilibrium aggregate obfuscation $\bar{\epsilon}^*$ in the dynamic game with a learning producer, as a function of the consumer discount factor β (3 seeds each, $T = 25000$). For all $\beta > 0$, $\bar{\epsilon}^* < 1$: obfuscation survives against a best-responding, belief-updating producer, and decreases with patience.

β	$\bar{\epsilon}^*$ (mean $\pm$ std)
0.00	0.936 ± 0.024
0.50	0.828 ± 0.042
0.90	0.689 ± 0.077
0.99	0.717 ± 0.054

patient consumers obfuscate more, since discounted future gain from corrupting the producer’s belief outweighs the immediate cost. The trend is monotone up to $\beta = 0.90$; the slight uptick at $\beta = 0.99$ lies within seed variation. Thus we reran the experiment with different seed settings.

Table 2: Equilibrium aggregate obfuscation $\bar{\epsilon}^*$ vs. consumer discount factor β (24 seeds, $T = 20000$). For all $\beta > 0$, $\bar{\epsilon}^* < 1$: obfuscation survives a best-responding, belief-updating producer, and decreases with patience, saturating near $\beta \approx 0.9$. Spearman over the raw $(\beta, \bar{\epsilon}^*)$ pairs ($n = 96$).

β	$\bar{\epsilon}^*$ (mean $\pm$ std)
0.00	0.914 ± 0.035
0.50	0.852 ± 0.048
0.90	0.740 ± 0.071
0.99	0.749 ± 0.059
Spearman $\rho = -0.769$, $p < 10^{-6}$	

With 24 seeds, $\beta = 0.99$ is statistically indistinguishable from $\beta = 0.9$: the trend decreases and then plateaus rather than reversing. The overall negative correlation is highly significant (Spearman $\rho = -0.769$ with $p < 10^{-6}$). The plateau suggests the obfuscation incentive saturates once consumers are sufficiently patient.